Strong-model reasoning with weak-model intervention

Threshold > 2 · k = 1 · One rollout per question · 4,571 questions

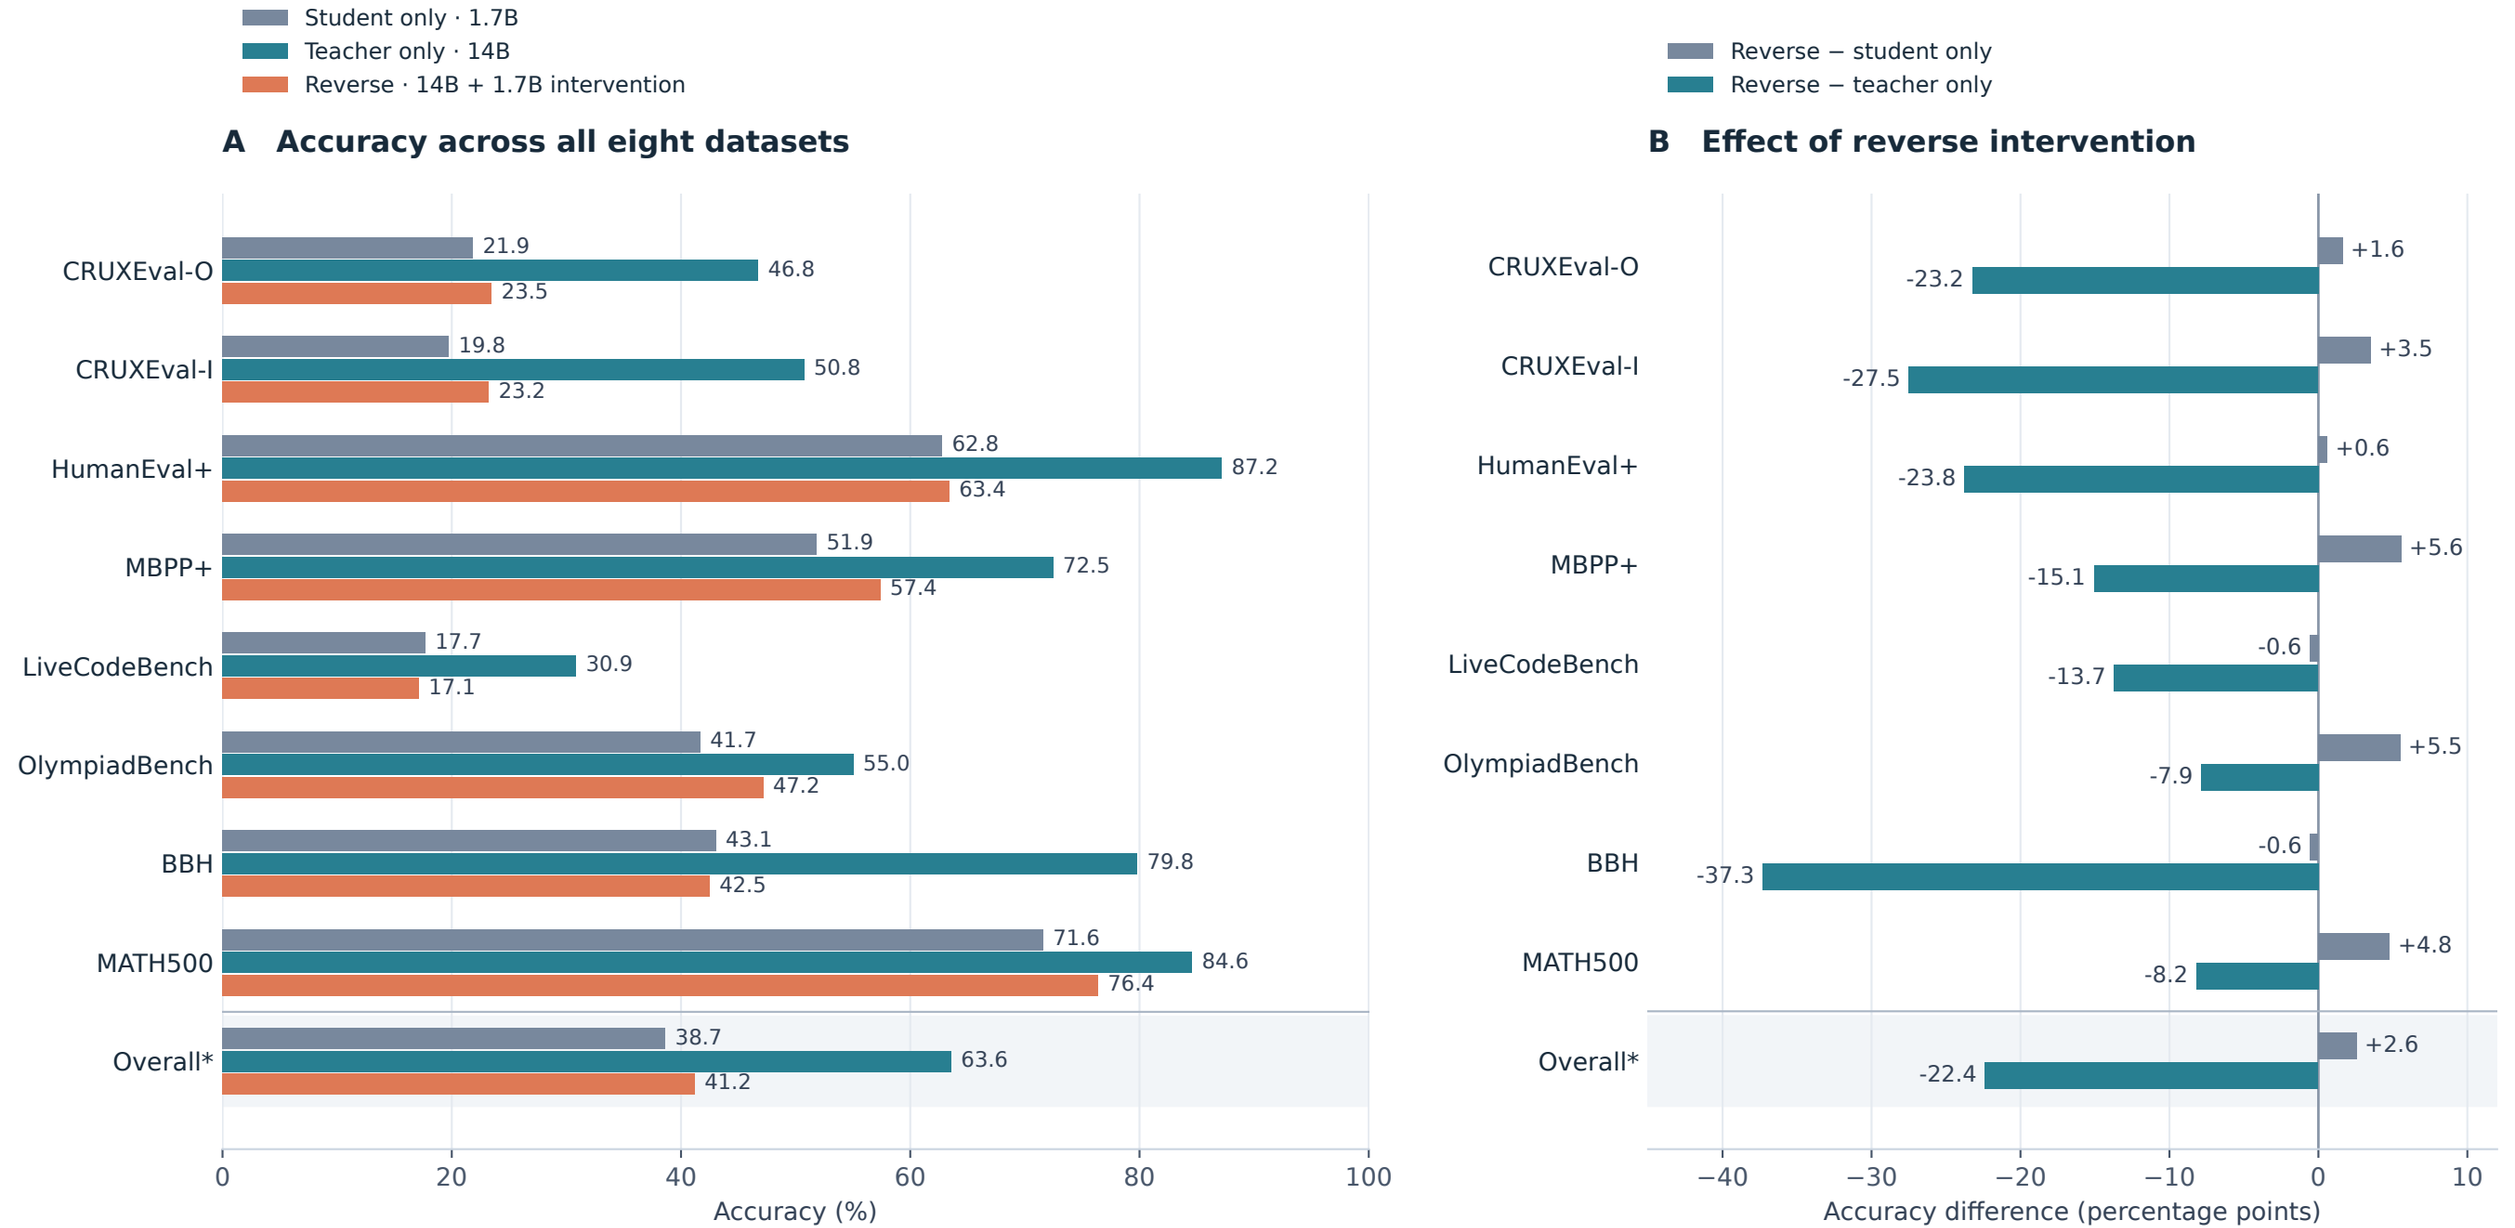

* Overall is weighted by question count. Generation timeouts count as incorrect (reverse: 244 / 4,571).

Historical single-rollout baselines; not an x8 mean. Reverse seed: 20260912; code baseline seed: 20260911.

Matched questions/prompts; same math judge prompt and code test harness. No uncertainty intervals estimated.